SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY —
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

Performance prediction of iterator based database systems with decision trees

Hoang-Nam Tran

SCHOOL OF COMPUTATION, INFORMATION AND TECHNOLOGY — INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

Performance prediction of iterator based database systems with decision trees

Performancevorhersage iteratorbasierter Datenbanksysteme mit Entscheidungsbäumen

Author:	Hoang-Nam Tran
Examiner:	Prof. Dr. Thomas Neumann
Supervisor:	Maximilian Rieger, M.Sc.
Submission Date:	27.04.2025

I confirm that this bachelor's thesis is my own work and I have documented all sources and material used.

Munich, 27.04.2025 Hoang-Nam Tran

Acknowledgments

Abstract

Contents

Acknowledgments **iv**

Abstract **v**

1 Introduction **1**
1.1 Section . . . 1
1.1.1 Subsection . . . 1

2 Theoretical Background **3**
2.1 Databases and Query Execution . . . 3
2.1.1 Iterator-Based Execution (PostgreSQL) . . . 3
2.1.2 Compiled Query Execution (Umbra) . . . 3
2.1.3 Pipelines and Pipeline Breakers . . . 3
2.2 Database Performance Prediction . . . 3
2.2.1 Benchmarking . . . 3
2.2.2 Error Metrics . . . 3
2.2.3 Traditional Cost Models and Learned Approaches . . . 3
2.2.4 Neural Networks, Decision Trees, and T3 . . . 3
2.3 Related Work . . . 3
2.4 ——————— . . . 3
2.5 Database Basics . . . 3
2.6 Performance Prediction . . . 3
2.7 Benchmarking . . . 4
2.8 Cost Models . . . 4
2.9 Decision Trees . . . 4
2.10 PostgreSQL . . . 4
2.11 Umbra . . . 4
2.12 T3 . . . 4
2.13 Section . . . 4
2.13.1 Subsection . . . 4

3 Methodology **6**
3.1 Problem Statement and Prediction Target . . . 6

3.2 Data and Preprocessing 6
3.3 Implementation 7
3.3.1 Mapping Parsed PostgreSQL Plans to the T3 Plan Format 7
3.3.2 Pipeline Decomposition 9
3.3.3 PostgreSQL-Native Feature Vectors 9
3.4 Training 11
3.4.1 Training Targets and Model Configuration 13
3.4.2 Data Splits and Reproducibility 13
3.5 Other Performance Prediction Implementations 13
3.6 Section 14
3.6.1 Subsection 14

4 Evaluation **16**
4.1 Ablation Study 16
4.1.1 PostgreSQL-native Features vs. Direct Transfer 16
4.1.2 Actual vs. Estimated Cardinalities 16
4.1.3 Pipeline vs Tuple targets 16
4.2 Section 16
4.2.1 Subsection 16

5 Future Work **18**
5.1 Section 18
5.1.1 Subsection 18

6 Conclusion **20**
6.1 Section 20
6.1.1 Subsection 20

Abbreviations **22**

List of Figures **23**

List of Tables **24**

Bibliography **25**

1 Introduction

1.1 Section

Citation test [Lam94].

Acronyms must be added in `main.tex` and are referenced using macros. The first occurrence is automatically replaced with the long version of the acronym, while all subsequent usages use the abbreviation.

E.g. `\ac{TUM}`, `\ac{TUM}` ⇒ Technical University of Munich (TUM), TUM

For more details, see the documentation of the `acronym` package[1].

1.1.1 Subsection

See Table 6.1, Figure 6.1, Figure 6.2, Figure 6.3.

Table 1.1: An example for a simple table.

A	B	C	D
1	2	1	2
2	3	2	3

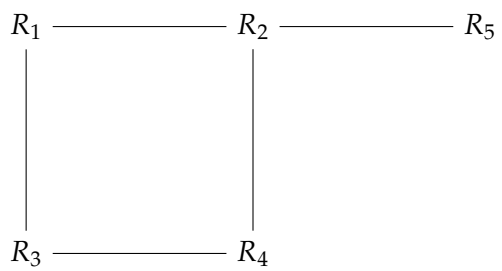

Figure 1.1: An example for a simple drawing.

[1] `https://ctan.org/pkg/acronym`

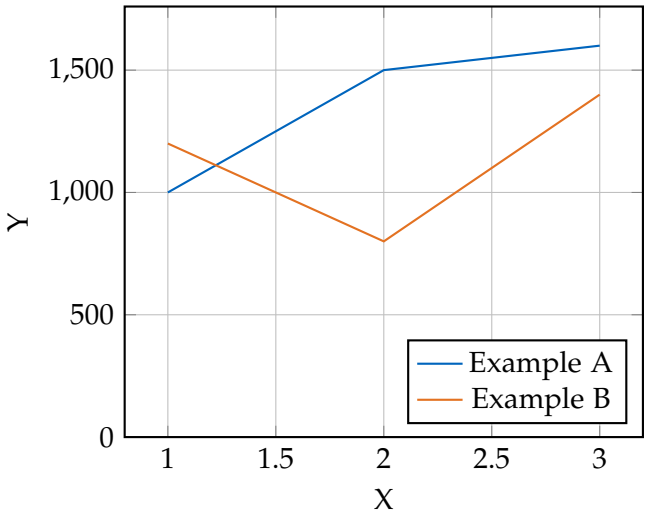

Figure 1.2: An example for a simple plot.

```
SELECT * FROM tbl WHERE tbl.str = "str"
```

Figure 1.3: An example for a source code listing.

2 Theoretical Background

2.1 Databases and Query Execution

2.1.1 Iterator-Based Execution (PostgreSQL)

2.1.2 Compiled Query Execution (Umbra)

2.1.3 Pipelines and Pipeline Breakers

2.2 Database Performance Prediction

2.2.1 Benchmarking

sowas auch schreiben wie ssb, tpch, etc.

2.2.2 Error Metrics

2.2.3 Traditional Cost Models and Learned Approaches

2.2.4 Neural Networks, Decision Trees, and T3

2.3 Related Work

2.4 ——————————

2.5 Database Basics

operator type, estimated/actual cardinalities, widths, and runtimes

2.6 Performance Prediction

draft: beschreiben, was performance prediction ist, was die ziele sind, etc. benchmarking, etc. Cost Model beschreiben

2.7 Benchmarking

timed queries, etc.

2.8 Cost Models

2.9 Decision Trees

2.10 PostgreSQL

2.11 Umbra

2.12 T3

2.13 Section

Citation test [Lam94].

Acronyms must be added in `main.tex` and are referenced using macros. The first occurrence is automatically replaced with the long version of the acronym, while all subsequent usages use the abbreviation.

E.g. `\ac{TUM}, \ac{TUM}` ⇒ TUM, TUM

For more details, see the documentation of the `acronym` package[1].

2.13.1 Subsection

See Table 6.1, Figure 6.1, Figure 6.2, Figure 6.3.

Table 2.1: An example for a simple table.

A	B	C	D
1	2	1	2
2	3	2	3

[1] `https://ctan.org/pkg/acronym`

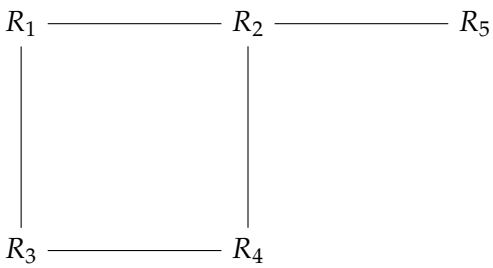

Figure 2.1: An example for a simple drawing.

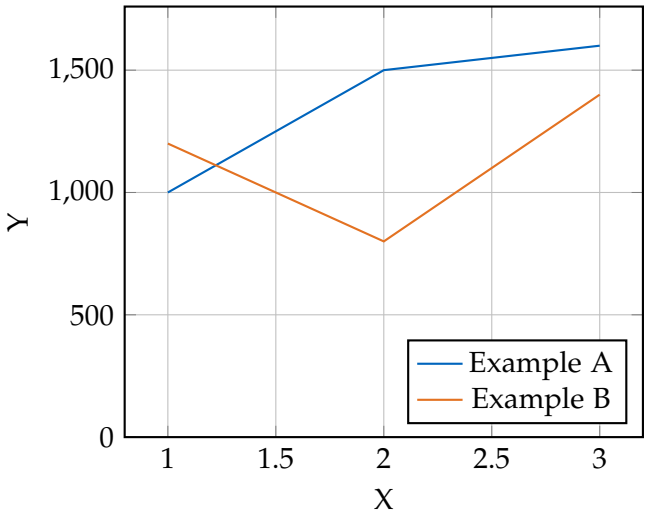

Figure 2.2: An example for a simple plot.

```
SELECT * FROM tbl WHERE tbl.str = "str"
```

Figure 2.3: An example for a source code listing.

3 Methodology

3.1 Problem Statement and Prediction Target

In this thesis, we address the transfer of the T3 runtime prediction approach from Umbra to PostgreSQL parsed plans. The central methodological goal is to preserve T3's core modeling principles. Upon that, we adapt plan conversion and feature extraction to PostgreSQL-specific plan information.

Also, the prediction target is total query runtime. Following T3, each query plan is decomposed into pipelines and the model predicts a duration for each pipeline. The predicted query runtime is the sum of these per-pipeline predictions.

3.2 Data and Preprocessing

Regarding the training data, we use the public Zero-Shot benchmark data introduced with the Zero-Shot Cost Model. The benchmark is particularly suitable for this thesis because it consists of real-world datasets (e.g., IMDB) and established benchmark workloads (e.g., SSB, TPC-H and JOB). Overall, the benchmark contains 20 databases with large variation in tables, columns, and foreign-key relationships [HB22].

Having this multitude of databases allows us to test the generalization capabilities of our model. Similar to the T3 paper, we can train the model on all but one databases and test on the remaining one. Moreover, as the widely used JOB workload is included, we can compare our performance predictions to other implementations later on. Further, PostgreSQL is the intended database system of our implementation, which avoids the effort of building a separate PostgreSQL benchmark.

Following the benchmark construction described in the paper, queries are generated per database and executed on a PostgreSQL DBMS (v12) on c8220 nodes on the CloudLab platform. Then, they are stored as raw traces (i. e., the output of `EXPLAIN ANALYZE`). Next, the authors preprocess the queries to remove very long-running queries (queries above 30 s are excluded) to keep the benchmark practical and reproducible [HB22]. A further preprocessing step is to parse the raw traces into structured plan representations. These `parsed_plans` (15000 per database) contain the operator nodes found in the raw traces, but organize them in a tree structure with node-wise attributes.

For every node, e.g., operator type, estimated/actual cardinalities, widths, and runtimes are stored in the `plan_parameters` field.

Figure 3.1 shows such a structured plan representation for a `Hash Join` and its children. On top of the already mentioned information, also filter predicates can be retrieved from the `filter_columns` field. Furthermore, PostgreSQL-specific fields like `workers_planned` (planned parallel workers for the node) are available.

The tree structure can be used to create pipelines and the node information will be used to train the model. In the upcoming sections, we will describe these key implementation details.

Consequently, the parsed plans offer a good starting point for our pipeline and we can begin without further preprocessing.

3.3 Implementation

This section describes how a single parsed plan from the Zero-Shot benchmark is turned into a T3-style representation including pipelines. Then, we derive fixed-length feature vectors for each pipeline. The procedure follows the pipeline idea of T3 [RN25], but also replaces Umbra-specific tuple-size and stage-based features with properties that PostgreSQL exposes in the `parsed_plans`. Once all the plans are converted, we can train the model.

3.3.1 Mapping Parsed PostgreSQL Plans to the T3 Plan Format

Each training example starts from one element of the `parsed_plans` array in the benchmark JSON files. An object in the array is a rooted tree: every node carries `plan_parameters` (operator name, estimated and actual cardinalities, tuple width, planner costs, optional filter predicates, timings, and further fields) and a list `children` of child subtrees.

In `zeroshot_to_t3.py`, we convert this tree recursively into the same nested operator structure that the original T3 implementation expects for Umbra plans [RN25]. PostgreSQL operator names are mapped to abstract operators (`tablescan`, `join`, `sort`, and `groupby` nodes). The PostgreSQL-specific temporary operators `hash` and `materialize` are mapped to the abstract operator `temp`.

Regarding joins, the child order is normalized so that it matches T3's pipeline rules: i. e., the PG-build subtree becomes the *left* child and the PG-probe subtree the *right* child.

Since the feature vector uses the PostgreSQL-native features and operators of the `parsed_plans`, we need to keep the original, not converted `plan_parameters` available for the feature extraction. Hence, the converted abstract operator names are only used

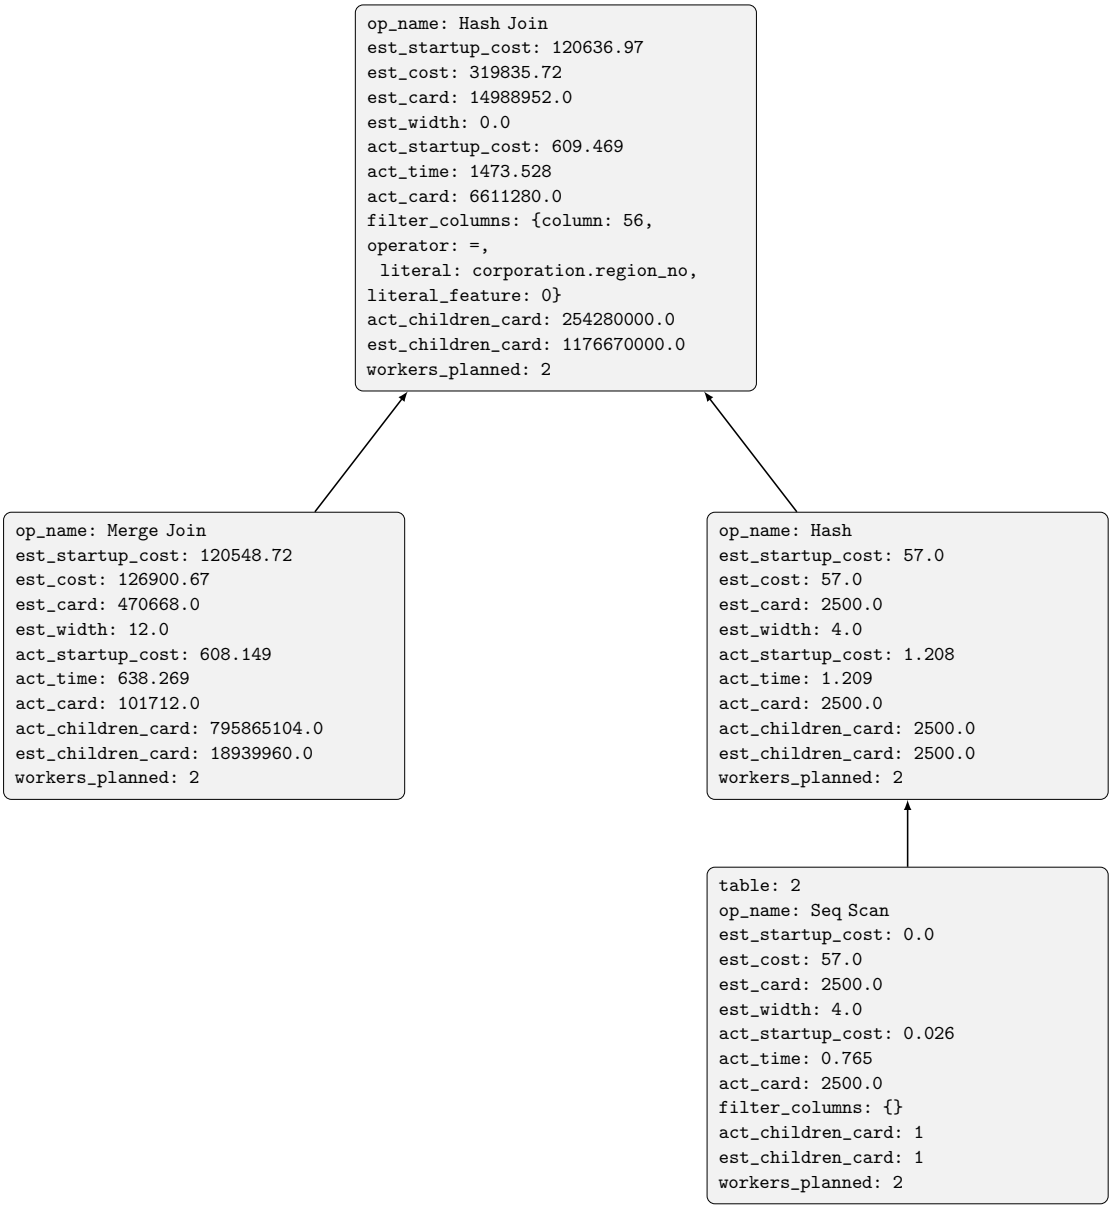

Figure 3.1: Top-node tree cutout of one `parsed_plan` with `Hash Join`. Child relationships are shown by arrows, and each node box contains selected `plan_parameters` of that node for the displayed cutout. Source: `credit/workload_100k_s1_c8220.json`, entry `0`.

to match the internal T3 representation for pipeline assignment, but not for the feature extraction.

Consequently, every converted node additionally stores a copy of the original `plan_parameters` under a dedicated `pg` payload. This payload also includes scan predicates which are stored in the `filter_columns` of the `plan_parameters`. Training and inference that use PostgreSQL-native features read from this payload and from the pipeline layout, which will be described in the next sections.

T3 also requires database schema information. However, the Zero-Shot data does not provide this information, so a minimal placeholder schema is used instead.

3.3.2 Pipeline Decomposition

Umbra plans group operators into *pipelines*: sets of operators that can produce tuples in a pipelined fashion, separated by *pipeline breakers*. Namely these are operators/stages, where data is materialized. For instance, sorting, grouping, or materialization, but also the build side of a hash join [RN25].

Since PostgreSQL plans do not provide such a pipeline breakdown, we need to derive it from the converted tree. We assign each node to a pipeline by walking the converted tree with the same high-level rules as in the T3 work. Thus, sort and aggregate nodes break pipelines; a materialize node is treated as a breaker. Joins place the build side in the preceding pipeline and keep the join together with the probe side in the next pipeline. The pipeline assignment is stored in the `analyzePlanPipelines` list of the converted plan.

For every pipeline we record the list of operator ids, and we derive start and stop times. They stem from the per-node actual timings in `plan_parameters` and are accumulated, so that pipeline durations are available for further processing during training.

3.3.3 PostgreSQL-Native Feature Vectors

The original T3 feature mapper for Umbra builds a feature vector from *execution phases* within each pipeline (e.g. scan, build, probe phases) and from cardinalities and tuple sizes attached to those phases [RN25]. As PostgreSQL plans contain different data, we use `pg_features.py` to construct the feature vector for each pipeline after the aforementioned pipeline decomposition. It becomes clear, why the `plan_parameters` are kept in the converted tree (`pg` payload entry) as we use them for the feature extraction which will be explained now.

Once the feature vector is constructed, we can train the model. Our feature vector differs from the original T3 feature vector in the following ways:

1. **PostgreSQL-native operator features.** Parsed PostgreSQL plans expose one `EXPLAIN` node per iterator operator. As no execution stages are available, we aggregate features per operator (fixed layout per operator type). Hence, we do not include the stage-based features of the original T3 feature mapper. This matches the iterator-based PostgreSQL approach and avoids inventing synthetic build/probe phases for every node.

2. **PostgreSQL planner cost estimates.** We include sums and extrema of PostgreSQL's `est_cost` and `est_startup_cost` from `plan_parameters` in the feature vector. These are optimizer-internal estimates of work and startup cost and are an additional PostgreSQL-specific signal that is not available in the Umbra feature vector. Consequently, we can include this information in the feature vector (and model training).

3. **Planned parallelism.** Another PostgreSQL-specific signal is the planned parallelism. So, we take the sum of `workers_planned` over all operators in the pipeline. Hereby, the model can see how much intra-node parallelism the PostgreSQL planner expected (zero on purely sequential plans).

4. **Sizes.** T3 uses tuple sizes in its feature vector. This information is provided by the Umbra plans. However, the parsed PostgreSQL do not contain a perfect equivalent. Instead, they report a single width estimate per node (`est_width`). Thus, an average of `est_width` over all operators in the pipeline is used to represent sizes in the feature vector, replacing the tuple size features of T3.

In the Umbra feature vector of T3, rows are indexed by both operator type *and* execution phase (scan, build, probe, ...) like `HashJoin_Probe...` [RN25]. Now with the PostgreSQL implementation, we miss out on the execution phase dimension.

Whereas not an equivalent substitute, we replace this information with a pipeline-level dimension and keep the operator-level features for the individual operators. The pipeline-level block strives to produce metadata that superceeds the operator level, e. g., by using averages and sums over operators.

Consequently, each pipeline's feature vector consists of two concatenated blocks, as Figure 3.2 shows:

1. **Pipeline-level aggregates.** From the available `pg` payload, we try to extract as distinct information as possible.

   Therefore, we start by summarizing different information categories per pipeline with sum, max, and average statistics. The `pg_act_card_...` fields capture cardinality summaries.

Insights about PostgreSQL-specific cost, parallelism, and width features are captured in the `pg_est_cost_...`, `pg_est_startup_...`, `pg_workers_planned_sum` and `pg_est_width_avg` fields.

The `pg_num_(scan/join/sort/etc.)` family counts operator classes.

The `pg_scan_has_filter` and `pg_filter_(and/or/compare/like/in)` counters enumerate scan predicate attributes.

Finally, pipeline-context features describe structure using `pg_pipeline_num_ops` for instance.

2. **Operator-level block.** For every operator that occurs in `parsed_plans`, we assign a slice in the feature vector. If multiple operators of the same type occur in the pipeline, we sum the features over all occurences. In figure 3.2, the excerpt of the operator-type block for the `Nested_Loop join` operator is depicted.

   Depending on the operator type, we include different features. This feature set is derived from T3's Umbra implementation. For example, a join would include different features than a scan, like in the Umbra/T3 approach.

   Regarding the calculations of the features, we follow the Umbra/T3 implementation where adequate. A deviation would be, e. g, how expression percentages are formed due to the different nature of PostgreSQL's filter predicates.

   Finally, we include the operator count "`const`", which exists for every operator type in both our and the Umbra/T3 implementation.

In total, the feature vector has 154 dimensions (24 pipeline-level features + 130 operator-type features). Empty features are kept as zeroes to fix the feature vector dimension.

Regarding the cardinalities, either the actual or the estimated cardinalities are used in the feature vector, depending on which one is selected in training. However, even if estimates are used, the row names are still called "actual" cardinalities. This caveat is a design decision to keep the row names completely uniform.

Once the feature vectors for all pipelines are constructed, we get the feature matrix for the whole plan. With this matrix, we proceed to train the model in the next section.

3.4 Training

Concerning the training, we keep T3's LightGBM decision tree setup [RN25]. Then, we use the newly constructed PostgreSQL-native feature matrix. A main training difference is that we train directly on pipeline durations as labels rather than using tuple-based targets.

Pipeline-level block (excerpt; first 12 of 24 features)	
`pg_act_card_sum`	6541.0
`pg_act_card_max`	6537.0
`pg_est_cost_sum`	16882.95
`pg_est_cost_max`	16877.3
`pg_est_startup_sum`	125.82
`pg_workers_planned_sum`	4.0
`pg_est_width_avg`	10.0
`pg_num_scan`	1.0
`pg_num_join`	1.0
`pg_scan_has_filter`	0.0
`pg_filter_compare_count`	0.0
`pg_pipeline_num_ops`	2.0
Operator-type block (excerpt; Nested_Loop slice; 130 entries in full)	
`Nested_Loop_const`	1.0
`Nested_Loop_in_percentage`	0.11473152822395594
`Nested_Loop_right_card`	98378.0
`Nested_Loop_out_percentage`	0.008260670032124828
(remaining operator-type entries omitted)	

Figure 3.2: Excerpts of the pipeline-level block (first twelve features) and of the `Nested_Loop` entries in the operator-type block, for one pipeline from query `index_workload_100k_s2_c8220_890` on the TPC-H database. The full vector has 154 dimensions; remaining pipeline and operator entries are omitted.

3.4.1 Training Targets and Model Configuration

Concerning the training labels, we use the observed pipeline duration. In T3 [RN25], the training target is the per-tuple-time instead, which is the pipeline duration divided by a pipeline scan cardinality. We will compare the per-tuple-time approach with the pipeline duration approach in the REFERENCEABLATIONSTUDY.

Regarding inference, our predicted query runtime is the sum of the predicted pipeline durations.

Our model uses gradient boosted decision trees as implemented in LightGBM [Ke+17]. The booster is trained with a mean absolute percentage error (MAPE) objective on the training rows. This is in line with the original T3 work [RN25].

Furthermore, we use the same training parameters as in the original T3 work [RN25]. Hence, the pipeline duration targets t are transformed as $-\log(t)$ to stabilize training. Also following T3, we cap the number of boosting rounds at 200 trees to mitigate overfitting.

3.4.2 Data Splits and Reproducibility

Generalization. We match the training intention of the Zero-Shot benchmark [HB22] and of T3's cross-instance evaluation [RN25]. Both aim to measure whether the trained model transfers to an unseen database instance. Thus, we use a leave-one-out split on the `parsed_plan` files: We specify one instance as a holdout and use its `parsed_plan` files as the test set. All other instances form the training data.

Concretely, `training_zeroshot_tpch_holdout.py` uses the `--holdout` argument to specify the holdout instance (default: `tpc_h`). Then, we perform a training for all test instances with `run_all_holdouts.py`.

Train and Validation within the Training Data. Inside the training data, complete queries (entire plans) are partitioned at random into a training fraction of 80 % and a validation fraction of 20 %. The default seed is 42.

Environment. Model training and evaluation for the results reported in this thesis were run on a personal laptop: an Apple MacBook Pro (16-inch, 2021) with the Apple M1 Pro processor (10 CPU cores, 16 GPU cores), 16 GB memory, and a 512 GB SSD, under macOS 15.7.3.

3.5 Other Performance Prediction Implementations

Add other implementations that you want to compare to in evaluation.

3.6 Section

Citation test [Lam94].

Acronyms must be added in `main.tex` and are referenced using macros. The first occurrence is automatically replaced with the long version of the acronym, while all subsequent usages use the abbreviation.

E.g. `\ac{TUM}, \ac{TUM}` ⇒ TUM, TUM

For more details, see the documentation of the `acronym` package[1].

3.6.1 Subsection

See Table 6.1, Figure 6.1, Figure 6.2, Figure 6.3.

Table 3.1: An example for a simple table.

A	B	C	D
1	2	1	2
2	3	2	3

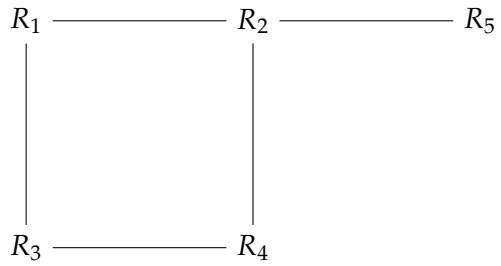

Figure 3.3: An example for a simple drawing.

[1] `https://ctan.org/pkg/acronym`

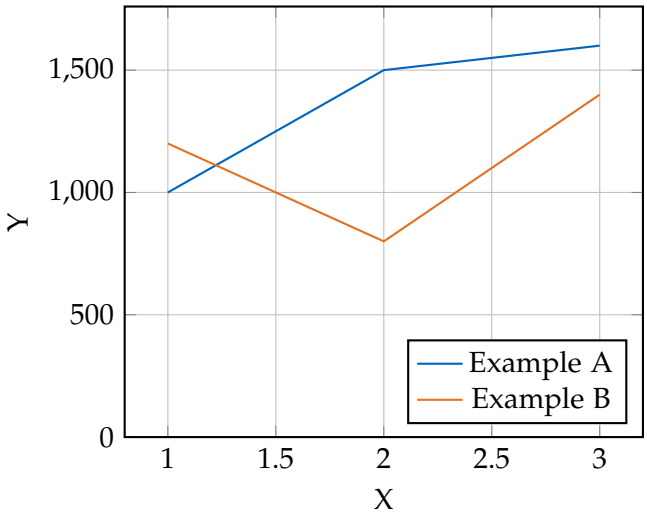

Figure 3.4: An example for a simple plot.

```
SELECT * FROM tbl WHERE tbl.str = "str"
```

Figure 3.5: An example for a source code listing.

4 Evaluation

4.1 Ablation Study

4.1.1 PostgreSQL-native Features vs. Direct Transfer

4.1.2 Actual vs. Estimated Cardinalities

4.1.3 Pipeline vs Tuple targets

4.2 Section

Citation test [Lam94].

Acronyms must be added in `main.tex` and are referenced using macros. The first occurrence is automatically replaced with the long version of the acronym, while all subsequent usages use the abbreviation.

E.g. `\ac{TUM}, \ac{TUM}` ⇒ TUM, TUM

For more details, see the documentation of the `acronym` package[1].

4.2.1 Subsection

See Table 6.1, Figure 6.1, Figure 6.2, Figure 6.3.

Table 4.1: An example for a simple table.

A	B	C	D
1	2	1	2
2	3	2	3

[1] `https://ctan.org/pkg/acronym`

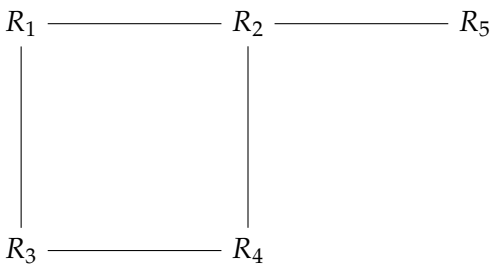

Figure 4.1: An example for a simple drawing.

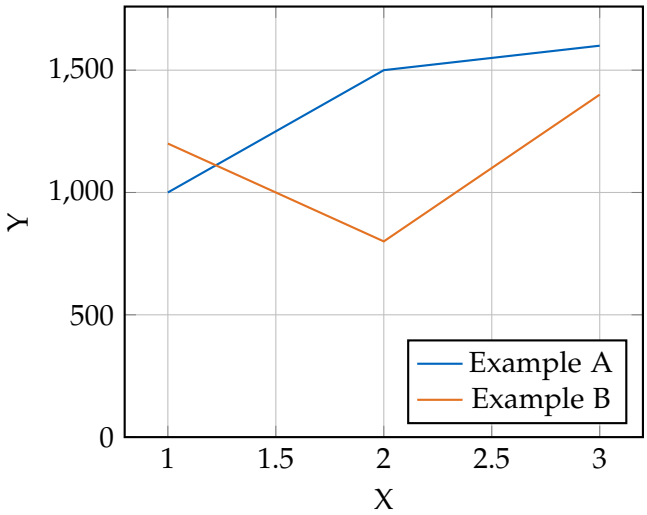

Figure 4.2: An example for a simple plot.

```
SELECT * FROM tbl WHERE tbl.str = "str"
```

Figure 4.3: An example for a source code listing.

5 Future Work

Data from Zeroshot kann limitierender faktor sein.

5.1 Section

Citation test [Lam94].

Acronyms must be added in `main.tex` and are referenced using macros. The first occurrence is automatically replaced with the long version of the acronym, while all subsequent usages use the abbreviation.

E.g. `\ac{TUM}, \ac{TUM}` ⇒ TUM, TUM

For more details, see the documentation of the `acronym` package[1].

5.1.1 Subsection

See Table 6.1, Figure 6.1, Figure 6.2, Figure 6.3.

Table 5.1: An example for a simple table.

A	B	C	D
1	2	1	2
2	3	2	3

[1] `https://ctan.org/pkg/acronym`

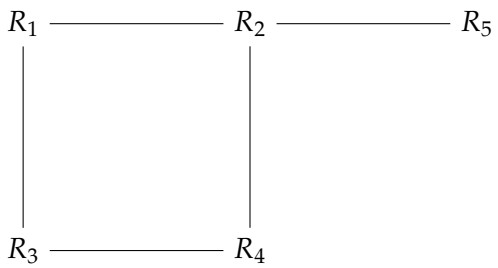

Figure 5.1: An example for a simple drawing.

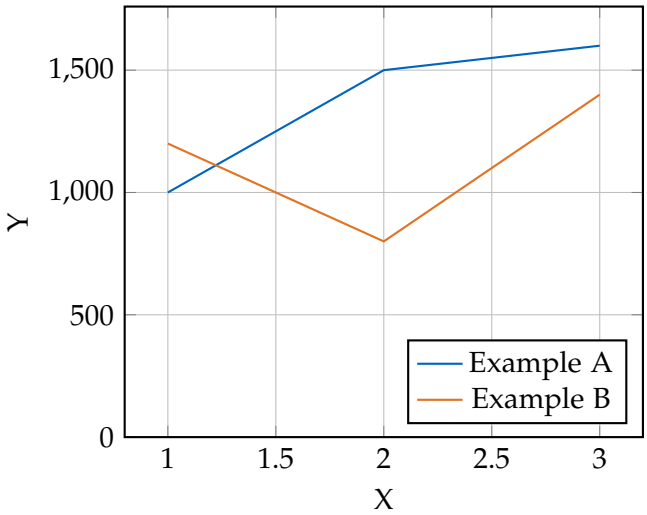

Figure 5.2: An example for a simple plot.

```
SELECT * FROM tbl WHERE tbl.str = "str"
```

Figure 5.3: An example for a source code listing.

6 Conclusion

6.1 Section

Citation test [Lam94].

Acronyms must be added in `main.tex` and are referenced using macros. The first occurrence is automatically replaced with the long version of the acronym, while all subsequent usages use the abbreviation.

E.g. `\ac{TUM}, \ac{TUM}` ⇒ TUM, TUM

For more details, see the documentation of the `acronym` package[1].

6.1.1 Subsection

See Table 6.1, Figure 6.1, Figure 6.2, Figure 6.3.

Table 6.1: An example for a simple table.

A	B	C	D
1	2	1	2
2	3	2	3

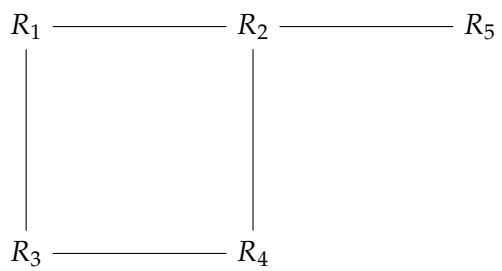

Figure 6.1: An example for a simple drawing.

[1] `https://ctan.org/pkg/acronym`

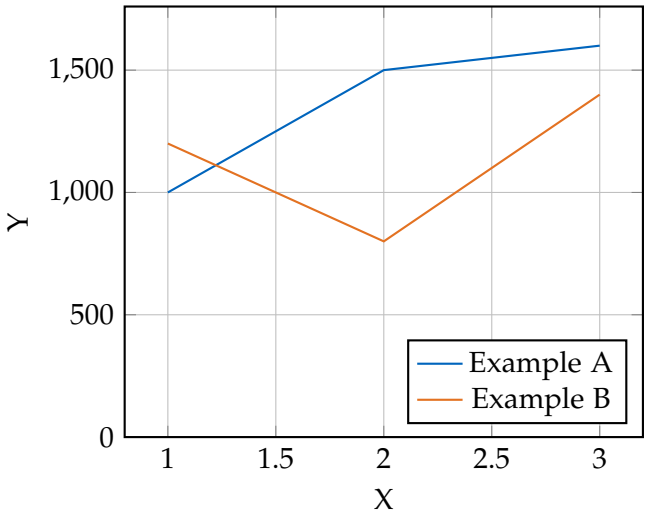

Figure 6.2: An example for a simple plot.

```
SELECT * FROM tbl WHERE tbl.str = "str"
```

Figure 6.3: An example for a source code listing.

Abbreviations

TUM Technical University of Munich

List of Figures

1.1 Example drawing 1
1.2 Example plot 2
1.3 Example listing 2

2.1 Example drawing 5
2.2 Example plot 5
2.3 Example listing 5

3.1 Top-node cutout with Hash Join 8
3.2 Excerpt of a PostgreSQL-native pipeline feature vector 12
3.3 Example drawing 14
3.4 Example plot 15
3.5 Example listing 15

4.1 Example drawing 17
4.2 Example plot 17
4.3 Example listing 17

5.1 Example drawing 19
5.2 Example plot 19
5.3 Example listing 19

6.1 Example drawing 20
6.2 Example plot 21
6.3 Example listing 21

List of Tables

1.1 Example table 1
2.1 Example table 4
3.1 Example table 14
4.1 Example table 16
5.1 Example table 18
6.1 Example table 20

Bibliography

[HB22] B. Hilprecht and C. Binnig. "Zero-shot cost models for out-of-the-box learned cost prediction." In: *Proc. VLDB Endow.* 15.11 (July 2022), pp. 2361–2374. ISSN: 2150-8097. DOI: 10.14778/3551793.3551799.

[Ke+17] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu. "LightGBM: A Highly Efficient Gradient Boosting Decision Tree." In: *Advances in Neural Information Processing Systems*. Vol. 30. 2017.

[Lam94] L. Lamport. *LaTeX : A Documentation Preparation System User's Guide and Reference Manual*. Addison-Wesley Professional, 1994.

[RN25] M. Rieger and T. Neumann. "T3: Accurate and Fast Performance Prediction for Relational Database Systems With Compiled Decision Trees." In: *Proc. ACM Manag. Data* 3.3 (June 2025). DOI: 10.1145/3725364.